

Bridging Performance and Affordability in AI Silicon

Juntaek Oh





Source: [조선일보](#)

엔비디아 GPU 26만장의 의미 ※B200 칩 기준

가격	연산 능력	HBM	필요 전력	국가 경쟁력
최대 104억달러 (약 15조원)	중대형 AI 데이터센터 5곳	HBM3E 208만개 탑재 (6000억~9000억원)	500~600MW (중형 화력발전소급)	미국, 중국에 이은 3위 수준

엔비디아가
국내에 공급하는
GPU

- 과학기술
정보통신부 5만장 국가 AI 컴퓨팅센터 구축·독자 AI 파운데이션 모델 개발
- 삼성전자 5만장 세계 최대 반도체 AI 팩토리 구축
- SK 그룹 5만장 제조 AI 클라우드 및 AI 팩토리 구축
- 현대차 그룹 5만장 AI 팩토리 통해 자율차, 스마트 팩토리, 로보틱스 강화
- 네이버 6만장 피지컬 AI 플랫폼 공동 개발

자료=업계 그림=이진명

Cassandra Unchained  @michaeljburry · Nov 10
Understating depreciation by extending useful life of assets artificially boosts earnings -one of the more common frauds of the modern era.

...
Massively ramping capex through purchase of Nvidia chips/servers on a 2-3 yr product cycle should not result in the extension of useful lives of compute equipment.

Yet this is exactly what all the hyperscalers have done. By my estimates they will understate depreciation by \$176 billion 2026-2028.

By 2028, ORCL will overstate earnings 26.9%, META by 20.8%, etc. But it gets worse. More detail coming November 25th. Stay tuned.

Network/Compute Depreciation Useful Life (Years)						
Company	2020	2021	2022	2023	2024	2025
META	3	4	4½	4½	4½	5½
GOOG	3	4	4	6	6	6
ORCL	5	5	5	5	6	6
MSFT	3	4	6	6	6	6
AMZN	4	4	5	5	6	5

Source: Company SEC Filings

한국경제  구독 PiCK ⓘ

"265조 사기극"...공매도 나선 '빅쇼트' 마이클 버리의 경고 [김인업의 퓨처 디스패치]

김인업 기자  입력 2025.11.11. 오전 10:26 · 수정 2025.11.11. 오전 11:02 기사원문

57 199 ...

버리 "AI데이터센터 감가상각 축소"
알파벳·MS 내용연수 3→6년 연장
이코노미스트 "시총 1100조 증발 우려"
엔비디아 GPU 개발 주기는 2→1년으로
오픈AI "주기 짧아지면 자금조달 어려워"
"AI 학습 아닌 추론 등 사용 가능" 주장도



사진=연합뉴스

AMD CEO Lisa Su says AI data center market will be worth \$1 trillion by 2030



AMD

팔로워 2,003,747명

5시간 ·

...

At AMD Financial Analyst Day, our CEO [Lisa Su](#) projected that the data center market will reach \$1 trillion by 2030, and AMD is positioned to lead that transformation.

“We have now all of the pieces to deliver full AI factories, and that is really our goal throughout this entire stack, across CPUs, GPUs, software, networking, and our cluster-level systems design.”

With leadership across AI infrastructure and multi-gigawatt deployments, AMD is powering the next era of high-performance computing.

Read more from [Yahoo Finance](#)

번역 표시



AMD CEO Lisa Su says AI data center market will be worth \$1 trillion by 2030

[finance.yahoo.com](#)

Meta plans \$600 billion US spend as AI data centers expand

By [Reuters](#)

November 7, 2025 10:23 AM PST · Updated November 7, 2025

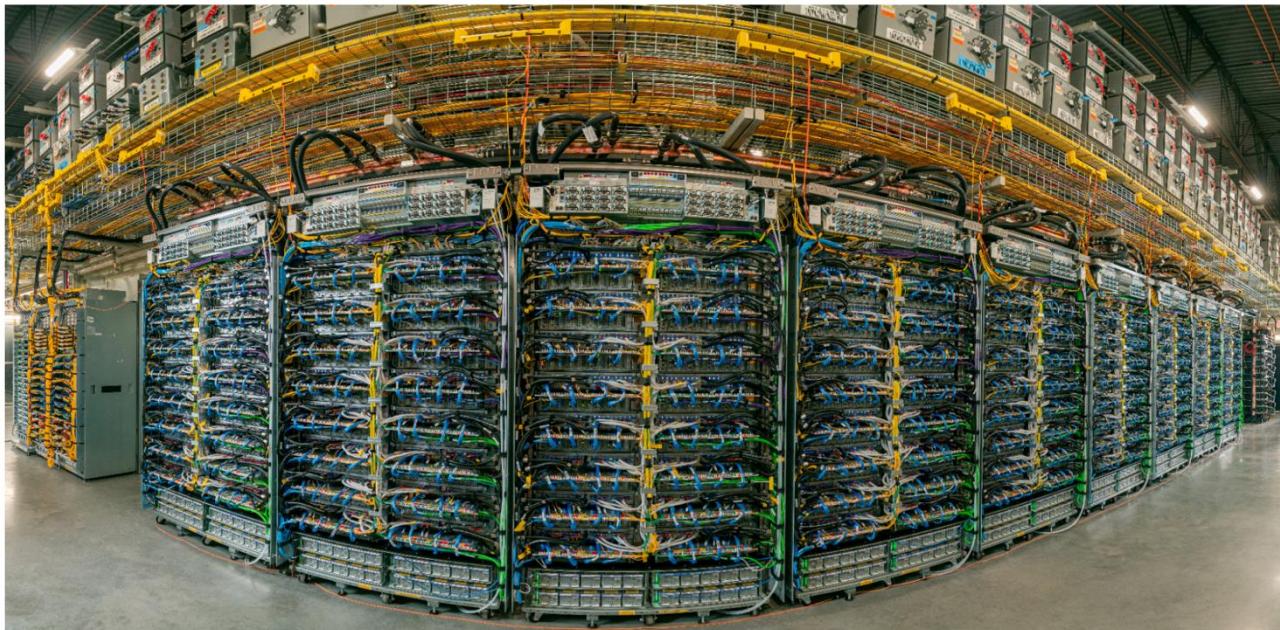


A teenager poses for a photo while holding a smartphone in front of a Meta logo in this illustration taken September 11, 2025.
REUTERS/Dado Ruvic/Illustration/File Photo [Purchase Licensing Rights](#)

Google Cloud Launches Ironwood TPUs, New Axion VMs for AI Inference



November 10, 2025
BY QUANTUM NEWS





cerebras raises

Groq Raises \$750 Million as Inference Demand Surges

The investment strengthens Groq's role in the American AI Stack, delivering fast, affordable compute worldwide.

LED BY



TIGERGLOBAL

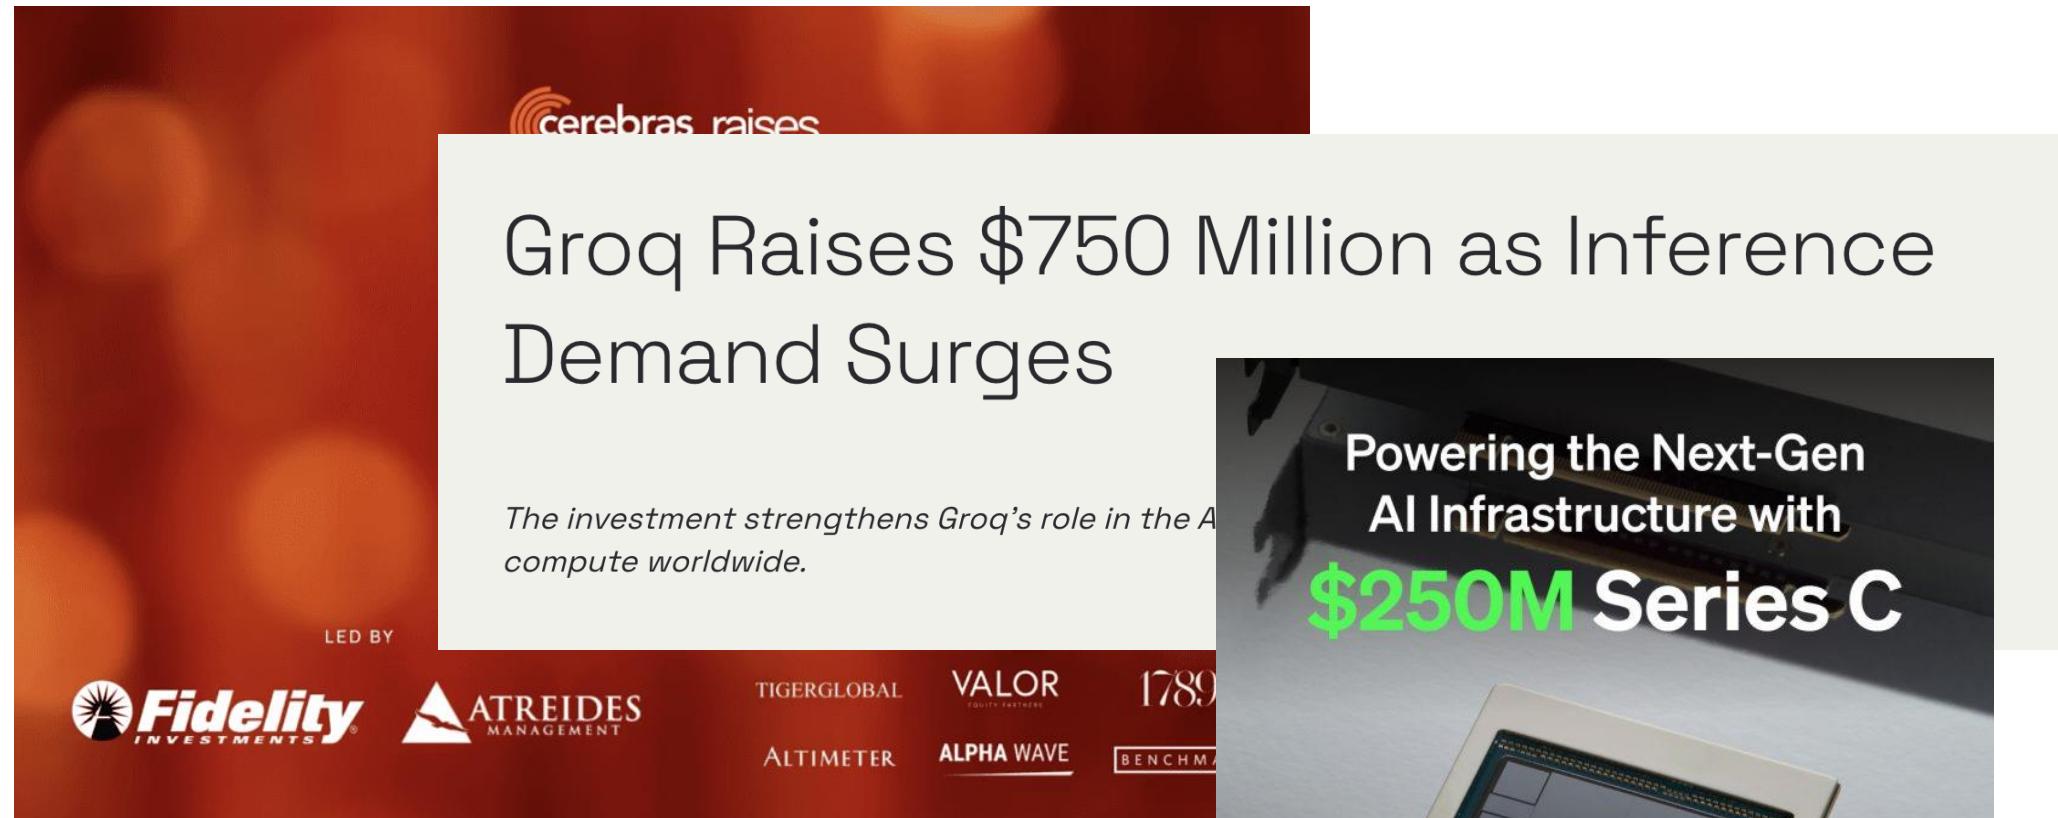
VALOR
EQUITY PARTNERS

1789
CAPITAL

ALTIMETER

ALPHA WAVE

BENCHMARK



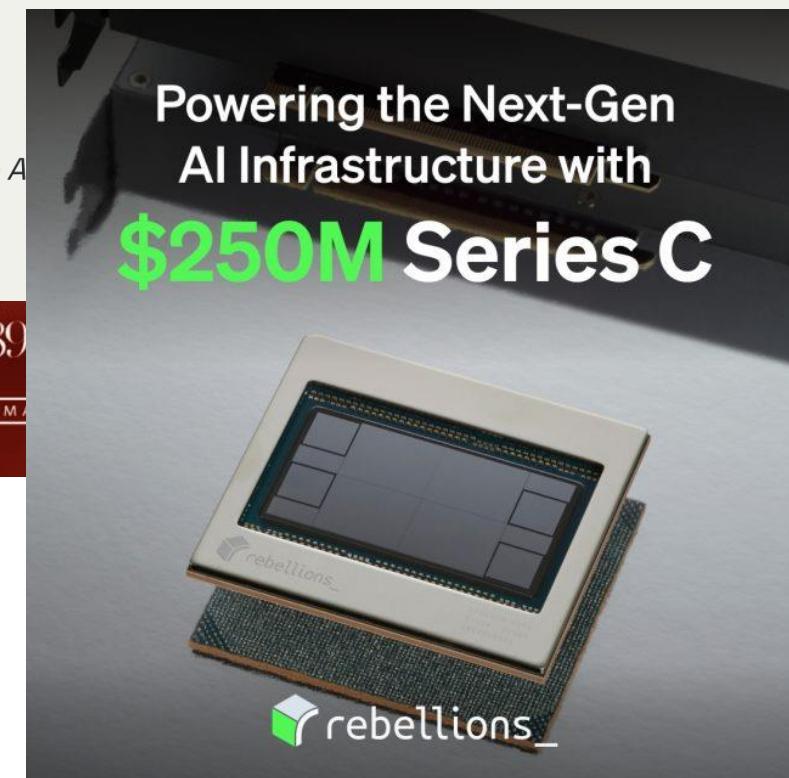
The banner features a red background with a blurred orange and yellow circular pattern. At the top left is the cerebras logo. The main title "Groq Raises \$750 Million as Inference Demand Surges" is centered in large, dark font. Below it is a subtitle: "The investment strengthens Groq's role in the AI compute worldwide." At the bottom, it says "LED BY" followed by logos for Fidelity Investments, ATREIDES MANAGEMENT, TIGERGLOBAL, VALOR EQUITY PARTNERS, 1789, ALTIMETER, ALPHA WAVE, and BENCHMARK.

Groq Raises \$750 Million as Inference Demand Surges

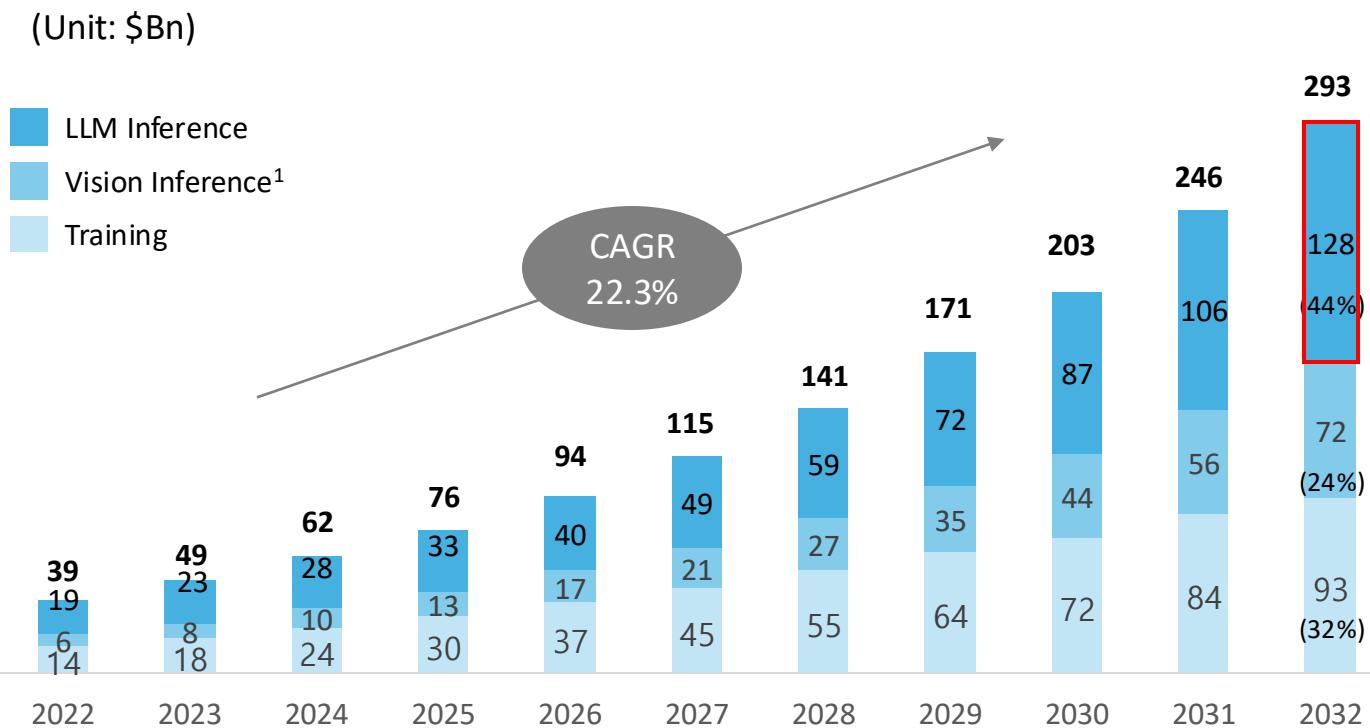
The investment strengthens Groq's role in the AI compute worldwide.

LED BY

Fidelity INVESTMENTS ATREIDES MANAGEMENT TIGERGLOBAL VALOR EQUITY PARTNERS 1789 ALTIMETER ALPHA WAVE BENCHMARK



AI Chips for LLM Inference are Poised to Lead the Market



LLM Inference Chip to Comprise 44% of the Total Market by 2032

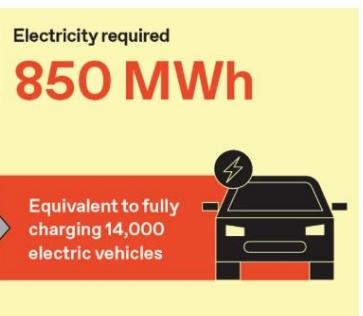
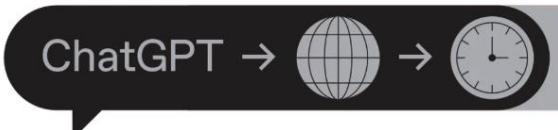
- Increasing demand for NLP applications across industries
- Growing complexity of language models
- Performance and efficiency benefits

2025

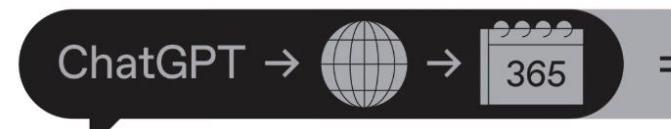
ChatGPT queries per average user per day
25



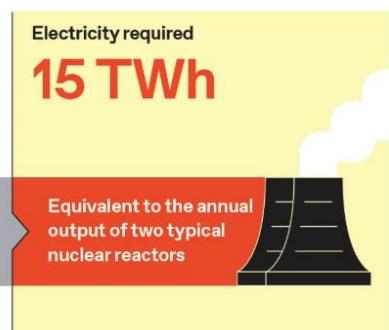
ChatGPT queries, all users per day
2,500,000,000



ChatGPT queries, all users per year
912,500,000,000



All generative AI queries, all users per year
5,100,000,000,000



2030

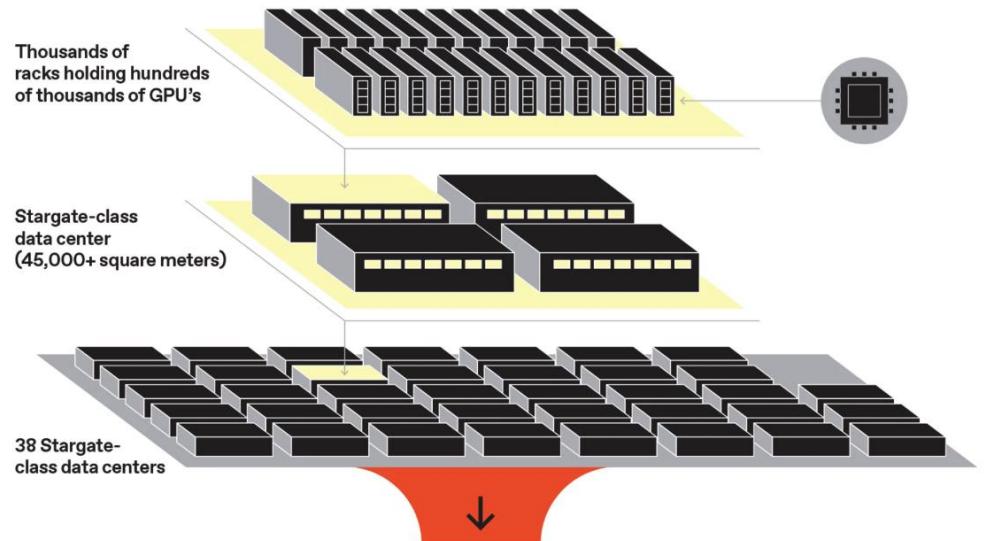
All generative AI queries, all users per year

120,000,000,000,000



Construction required

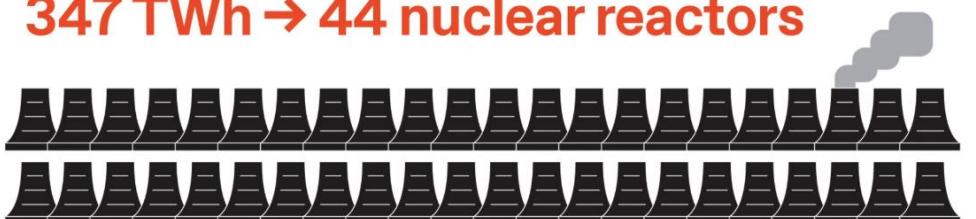
38 Stargate-class data centers



The Schneider Electric report estimates that all generative AI queries consume 15 TWh in 2025 and will use 347 TWh by 2030; that leaves 332 TWh of energy—and compute power—that will need to come online to support AI growth. That implies the construction of dozens of data centers along the lines of the Stargate Project, which plans to build the first ever 1-gigawatt facilities. Each of these facilities will theoretically consume 8.76 TWh per year—so 38 of these new campuses will account for the 332 TWh of new energy required.

Electricity required

347 TWh → 44 nuclear reactors



Why Microsoft's AI Chip Supply Chain Awaits Power

Explained: Generative AI's environmental impact
Rapid development and deployment of powerful generative AI models comes with environmental consequences, including increased electricity demand and water consumption.

Adam Zewe | MIT News
January 17, 2025



100MW data center capacity may remain unused for years due to power shortage

HyperAccel Steps Forward With Unique GenAI Chip Solutions to Win the AI Opportunity

Fast, Efficient, and Affordable for Generative AI Inference

HyperAccel is an GenAI Chip Startup hyperfocused on developing disruptive technologies and solutions specifically for Generative AI Inferencing. Designed from foundational understanding of transformer architecture Large Language Models, our LPU(LLM Processing Unit) aims for higher throughput performance with order of magnitude gains in cost and energy efficiency as compared to GPU's in market today. We dedicate ourselves to help our customers scale Fast, Efficient, and Affordable GenAI Services for Everyone!

Engineered for GenAI Leadership

¹ Llama 3.1 70B; 4096 Input/Output Token Length; Batch 32

HyperAccel Vision : GenAI is for Everyone

Specialized in LLM

We offer unique GenAI chips designed from first principles for LLM Inference – the HyperAccel LPU

Our ASIC models achieve **2x** higher throughput, **20x** better cost efficiency, and **5x** better energy efficiency compared to NVIDIA H100¹

Deep-Tech Founders & World-Class Team

Launched in January 2023, HyperAccel has grown to Series A and curated a team of over 70 world-class semiconductor HW, SW, and Systems professionals from leading companies

Disruptive Technologies

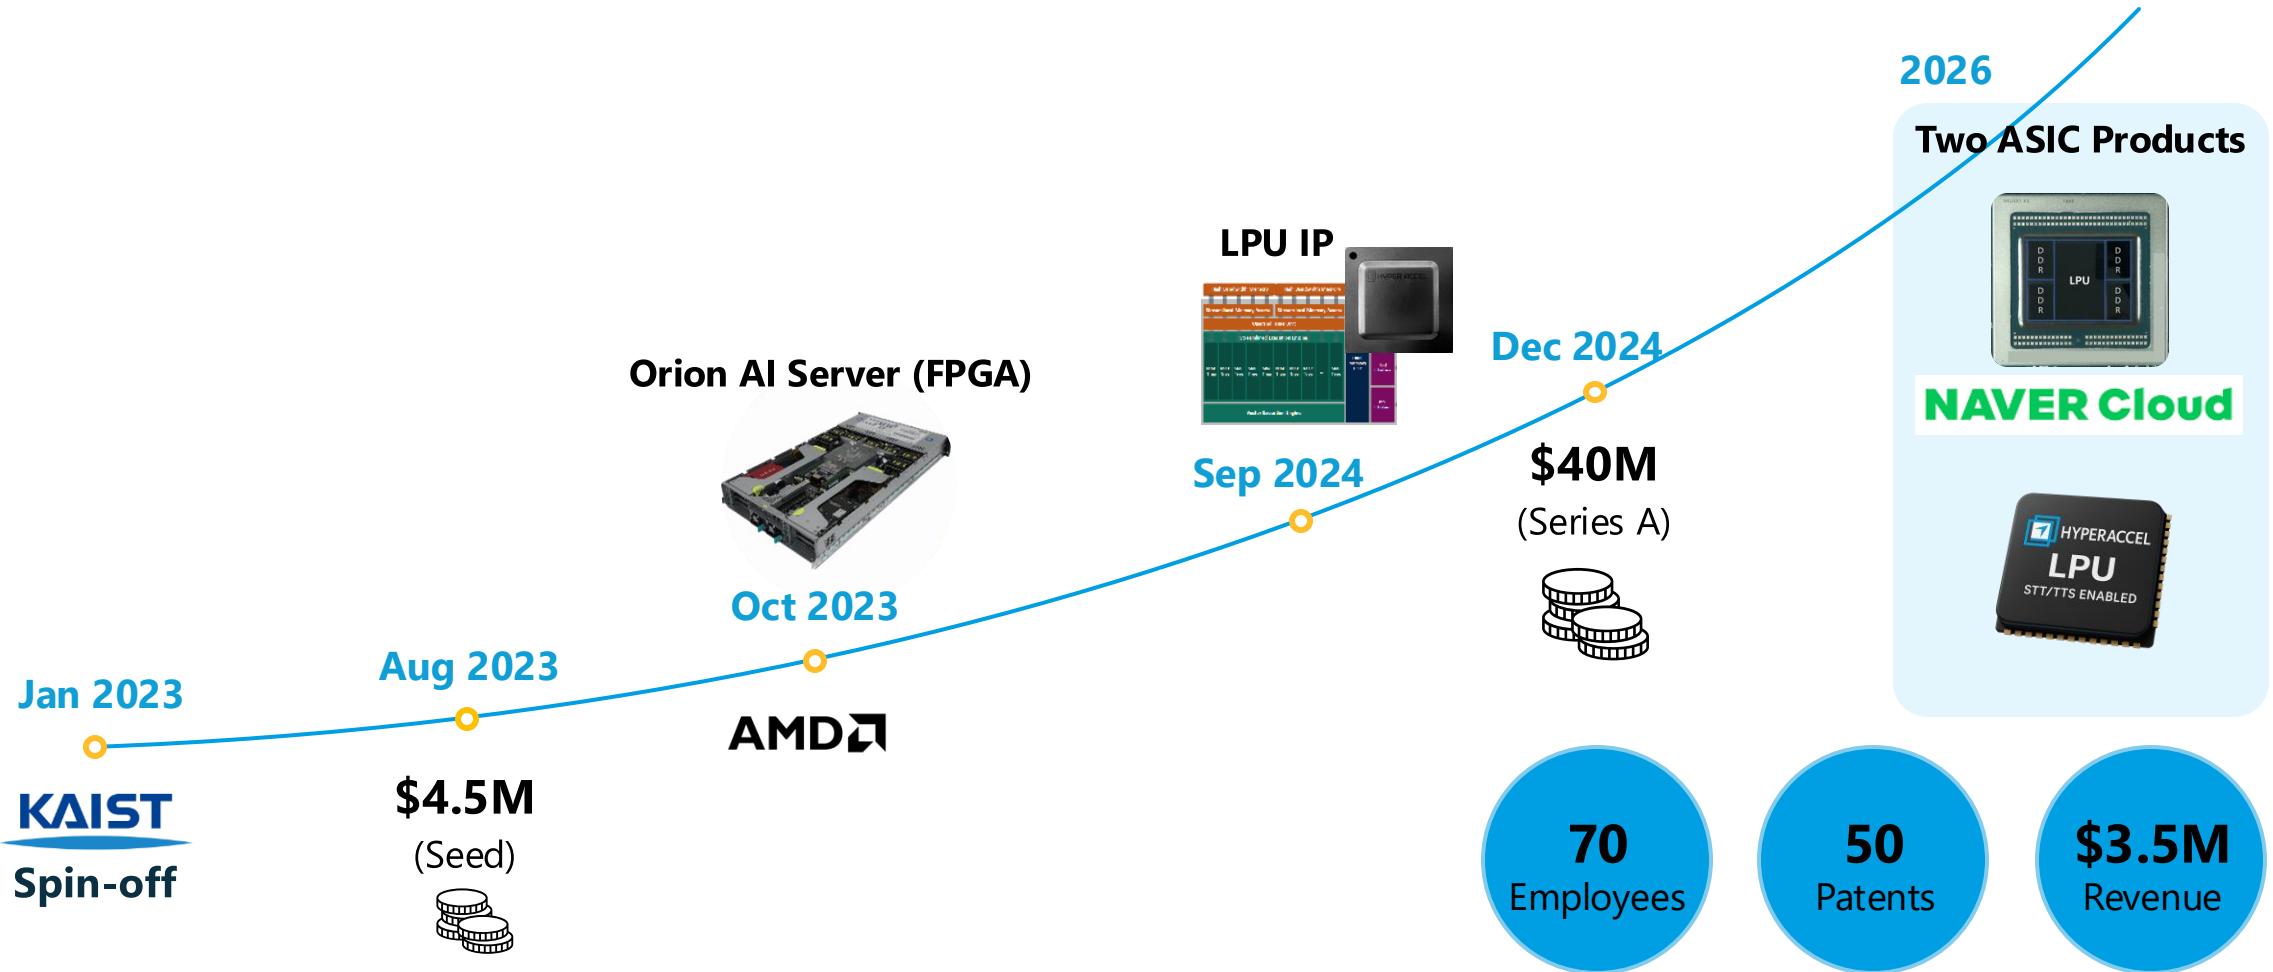
- Patented Hardware Designs
- Proprietary Full-Stack Software Platform
- Distributed/Runtime-Dynamic Execution
- Scalable/Extensible Architecture

Lighthouse Customers

NAVER Cloud for Datacenter AI ASIC

Global Consumer Electronics Company (TBA) for Edge AI ASIC

Company Profile



The Leadership Team

AI Computing & Infra

18 yrs Semiconductor Expert, 9 yrs @ Microsoft



Founder & CEO

Professor EE, KAIST
Head of AI Semiconductor Systems Research Lab

Engineering Leader, Microsoft Azure
2017 - 2019

Senior Researcher, Microsoft Research
2014 - 2017

Researcher, Microsoft Research
2012 - 2014

**Joo-Young
Kim**



World-Class Engineering Team

HW Engineering (24)

SW Engineering (25)

Systems Engineering (7)

NPU Computing

13 yrs Semiconductor Expert, 10 yrs @ Samsung



CTO

Neubla CTO
2021 - 2023

Samsung C-Lab Leader
2021 - 2021

Staff Engineer, Samsung S. LSI
2017 - 2020

SoC Engineer, Samsung S. LSI
2011 - 2015

**Jinwon
Lee**



Strategy & Business Development

20 yrs Management Expert, 13 yrs Tier-1 Consulting



CSO

OKIT Inc. Founder & CEO
2020 - 2024

Loplat COO
2017 - 2020

BCG Principal
2012 - 2017

Oliver Wyman Manager
2005 - 2011

**Yongwoong
Jung**



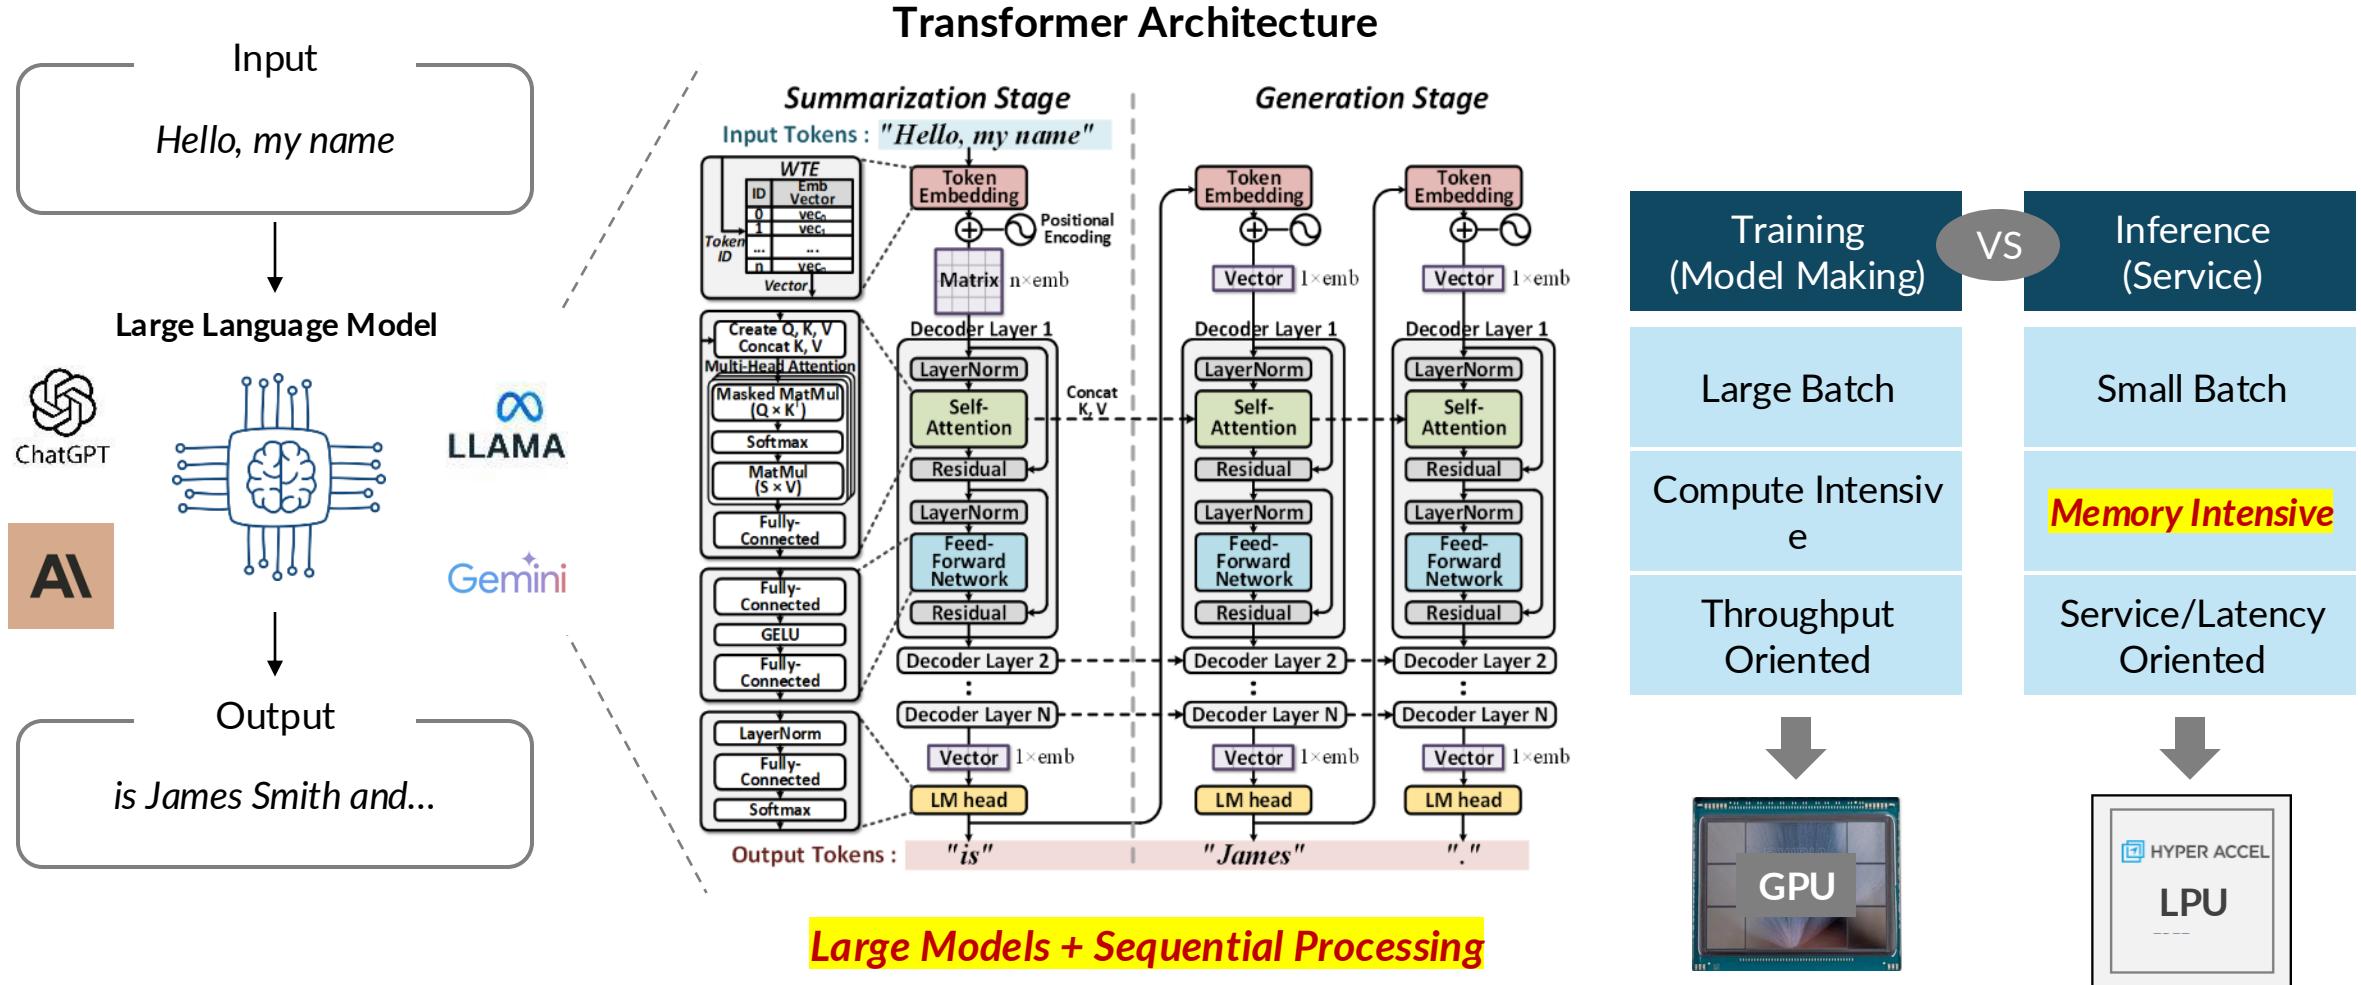
Education

INSEAD MBA (Class of 2012)
B.S., Korea University (2004)

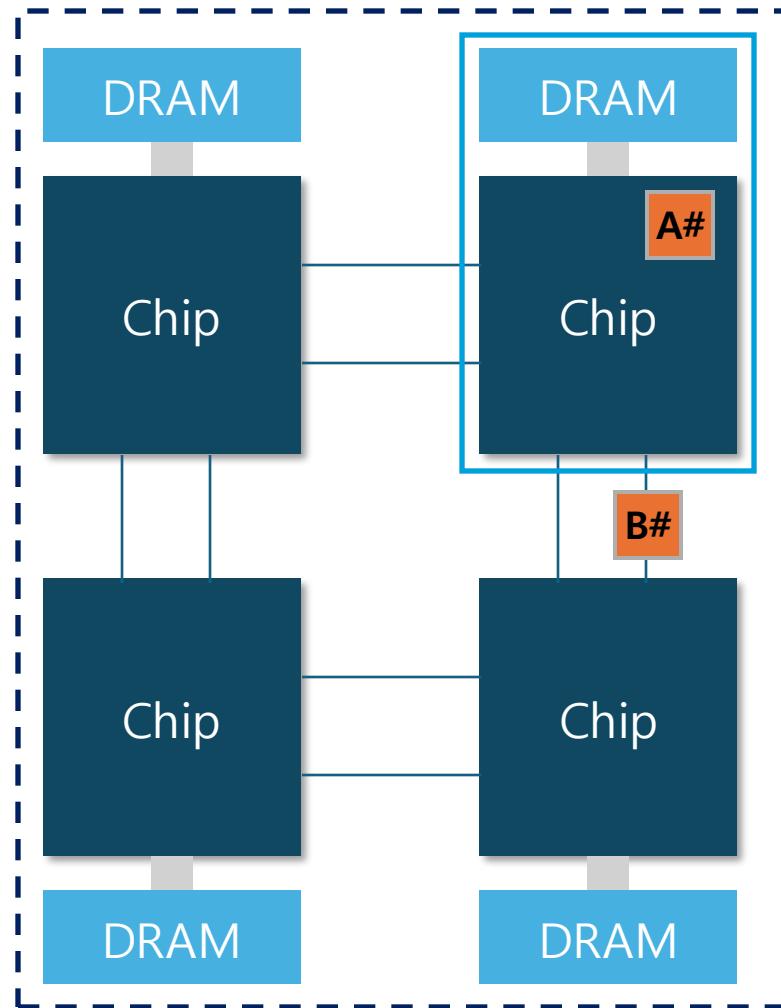
Business & Operations Team

BD, Sales, Mktg, Finance, HR, Admin
(8)

We Saw the Need to Specialize Having Deep Knowledge of LLM's the Effects of Large Models, Sequential Processing, and Memory Intensive for GenAI Inference



We Built Our LPU with Unique HW, SW, and System Designs to Deliver on Our Promise of High Performance and Affordability/Sustainability

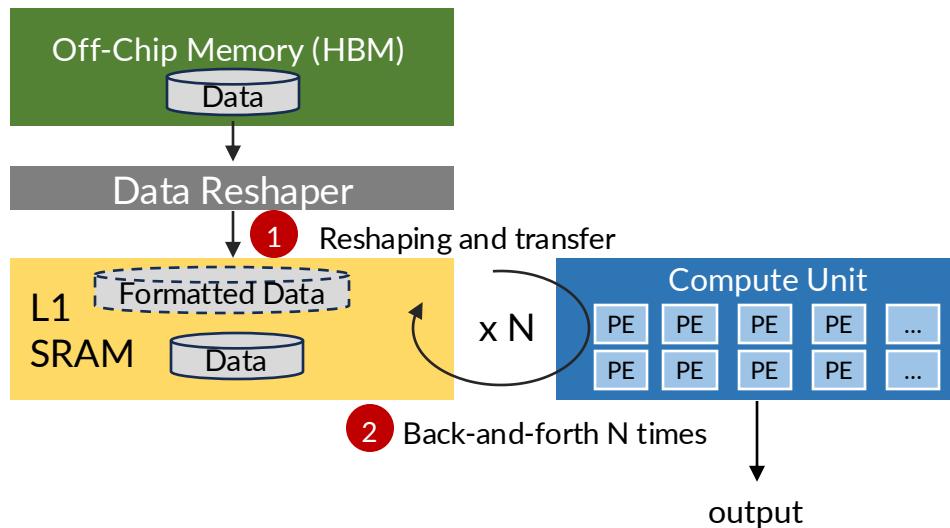


High Performance	Affordability/Sustainability
<ul style="list-style-type: none"> A1 Maximized Memory Bandwidth Utilization <ul style="list-style-type: none"> - Streamlined Memory Access A2 Specialized Compute Engines for End-to-End LLM Operations <ul style="list-style-type: none"> - Matrix and Vector Execution Engines A3 Adaptable HW Architecture Design for Chip-Level Scale Up/Down 	<p>Standard LPDDR Memory Replaces HBM</p> <ul style="list-style-type: none"> - Widely Available and Low Cost - 2x the Max Capacity of HBM - $\frac{1}{2}$ the Price of HBM - 60% less power consumption - No HBM Advanced Packaging Costs - Comparatively Lower IP Costs
<ul style="list-style-type: none"> B1 System-Level Extensibility via Peer-to-Peer Expandable Synchronization Link (ESL) B2 Full-Stack HyperDex Software Platform to Optimize the Performance of Multi-LPU System 	<p>Standard PCIe and Ethernet for NVLink-type Capabilities without Proprietary Charges and High Cost</p>

Streamlined Dataflow Architecture

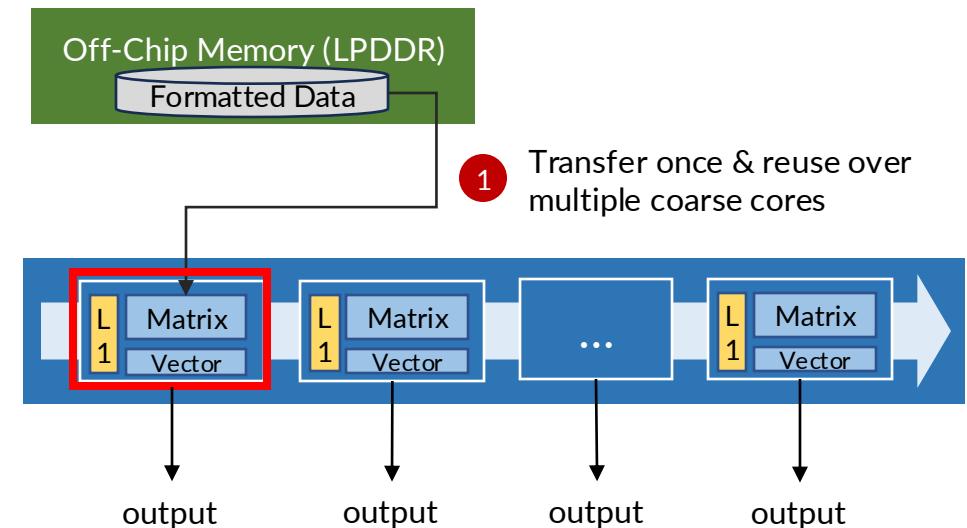
Conventional (GPU)

- 1000s of Small Cores
- Hierarchical access w/ lots of data movements



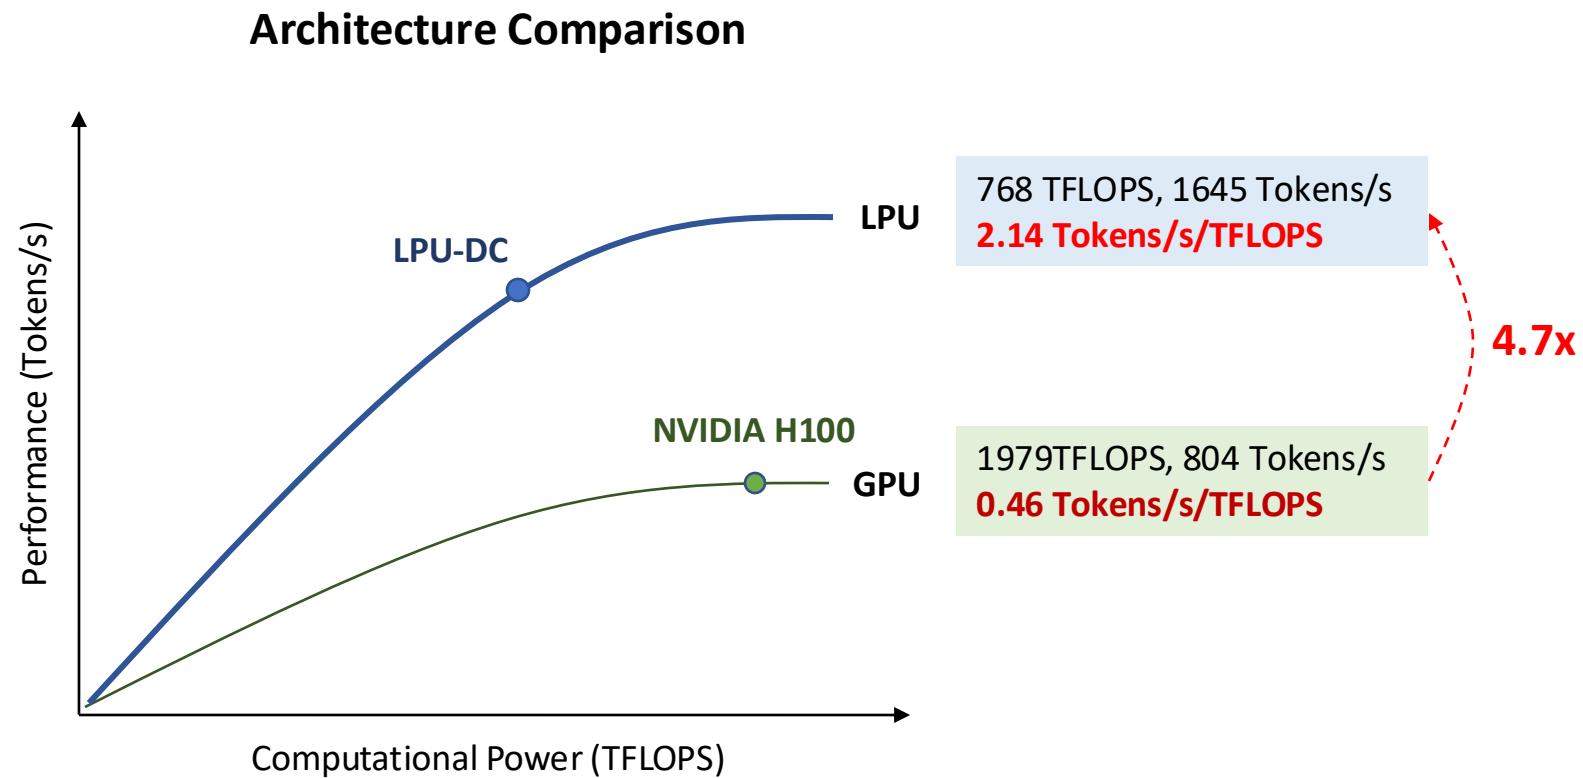
Ours (LPU)

- 10s of Big Cores
- Streamlined dataflow w/ maximum data reuse

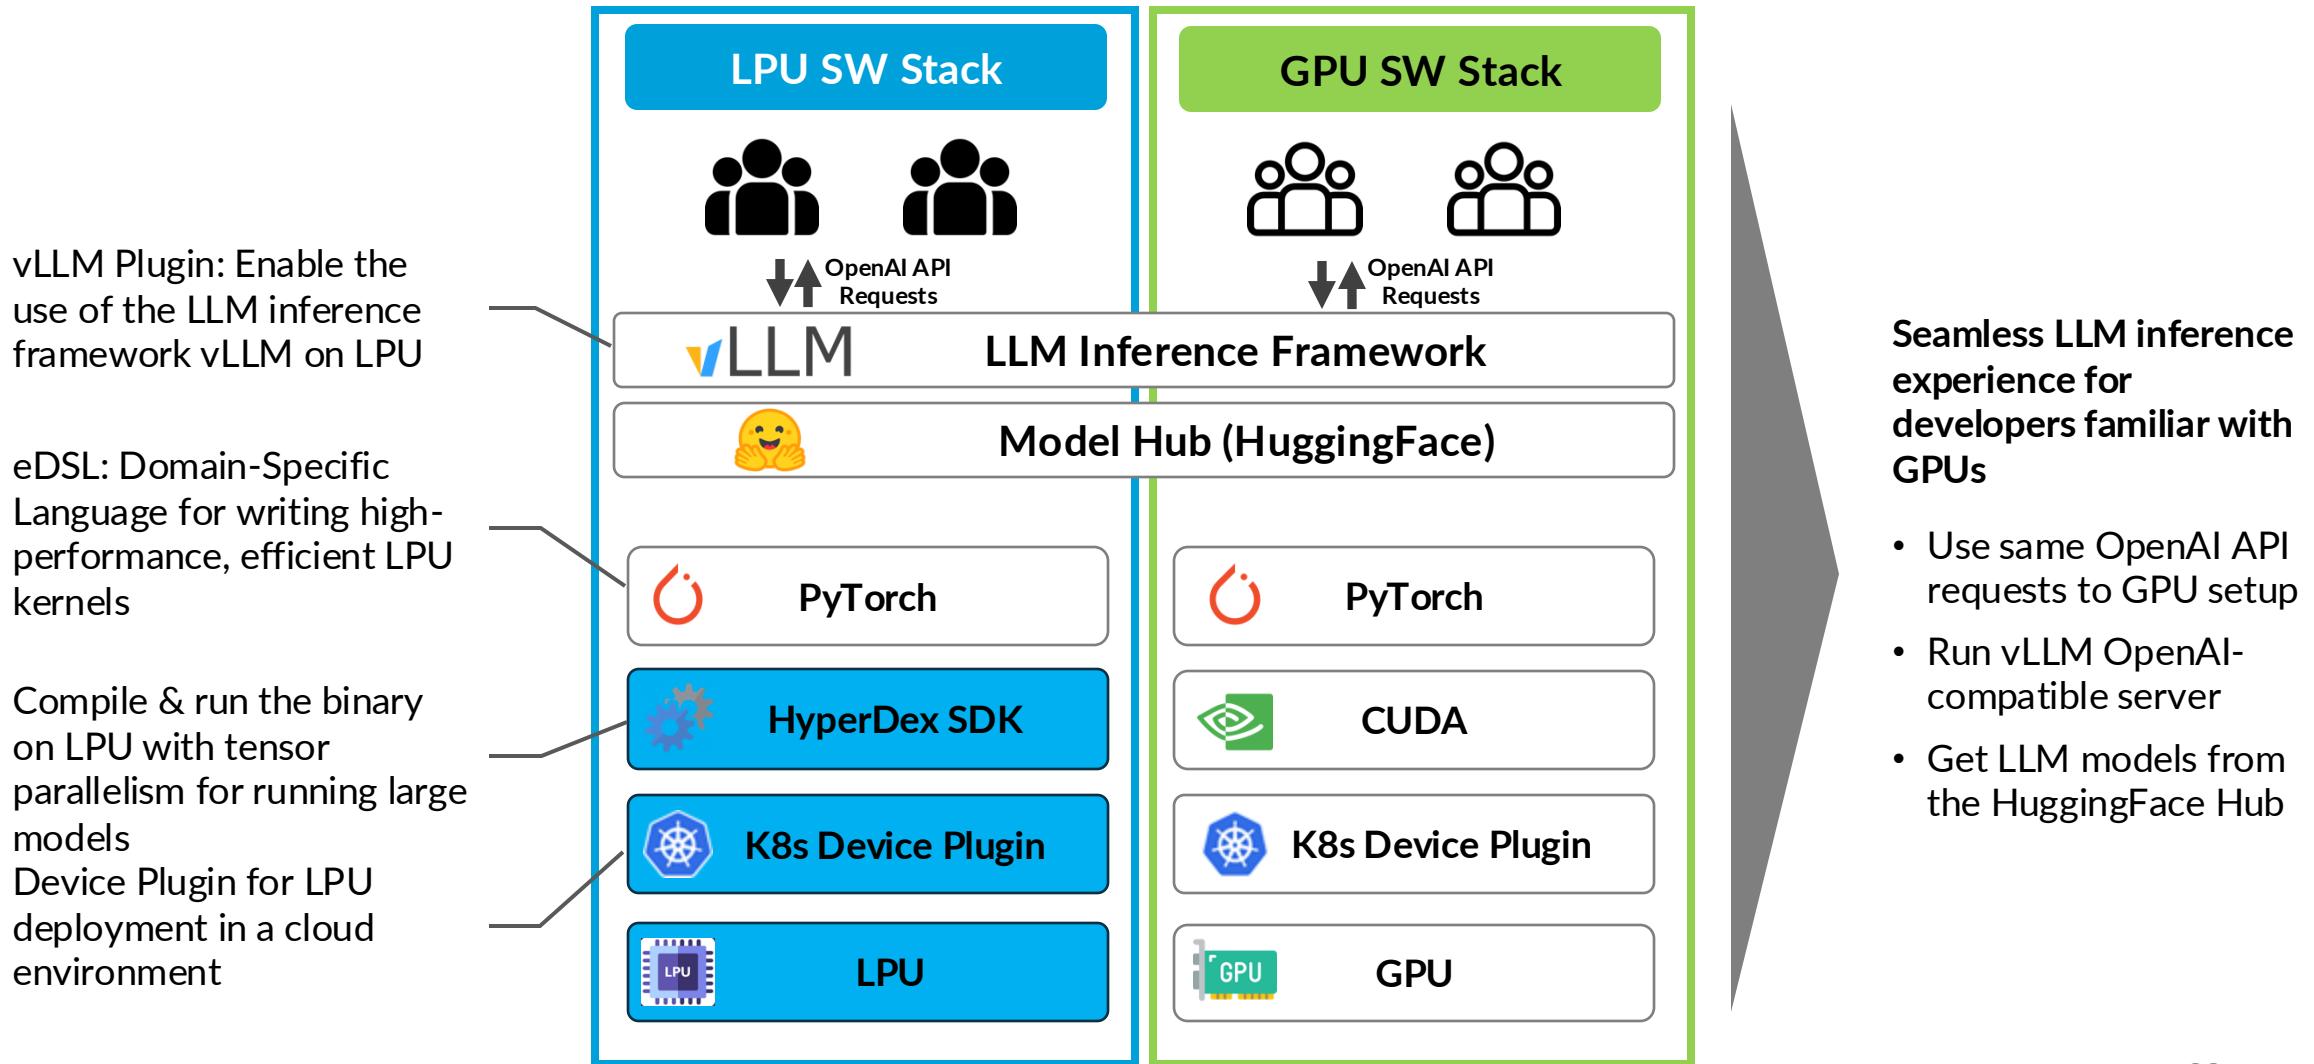


Maintain ~90% Utilization of Given Peak BW

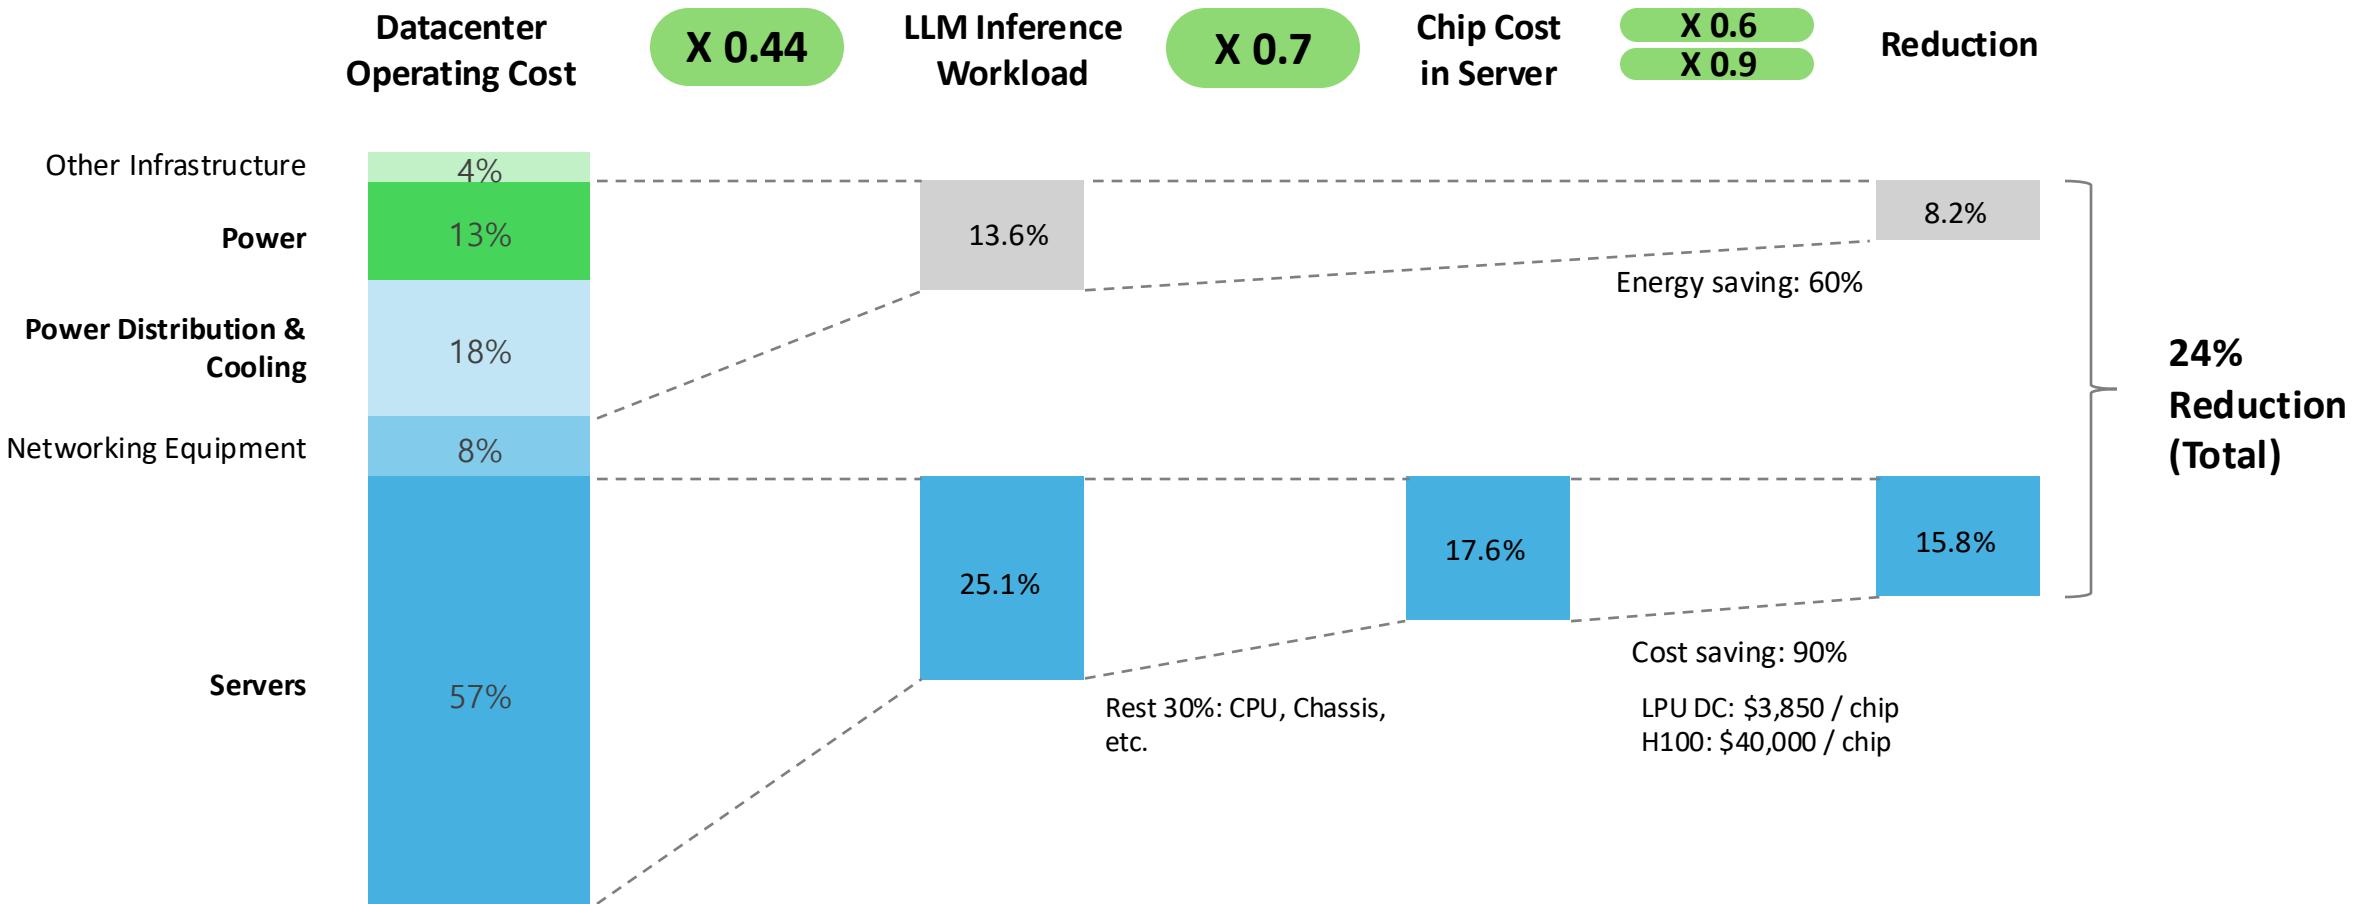
Achieving 5x LLM Performance per TFLOPS



Prioritizing Software, We Built From the Start Our Full-Stack HyperDex Software Platform to Optimize HW Performance and Engage GenAI Developers

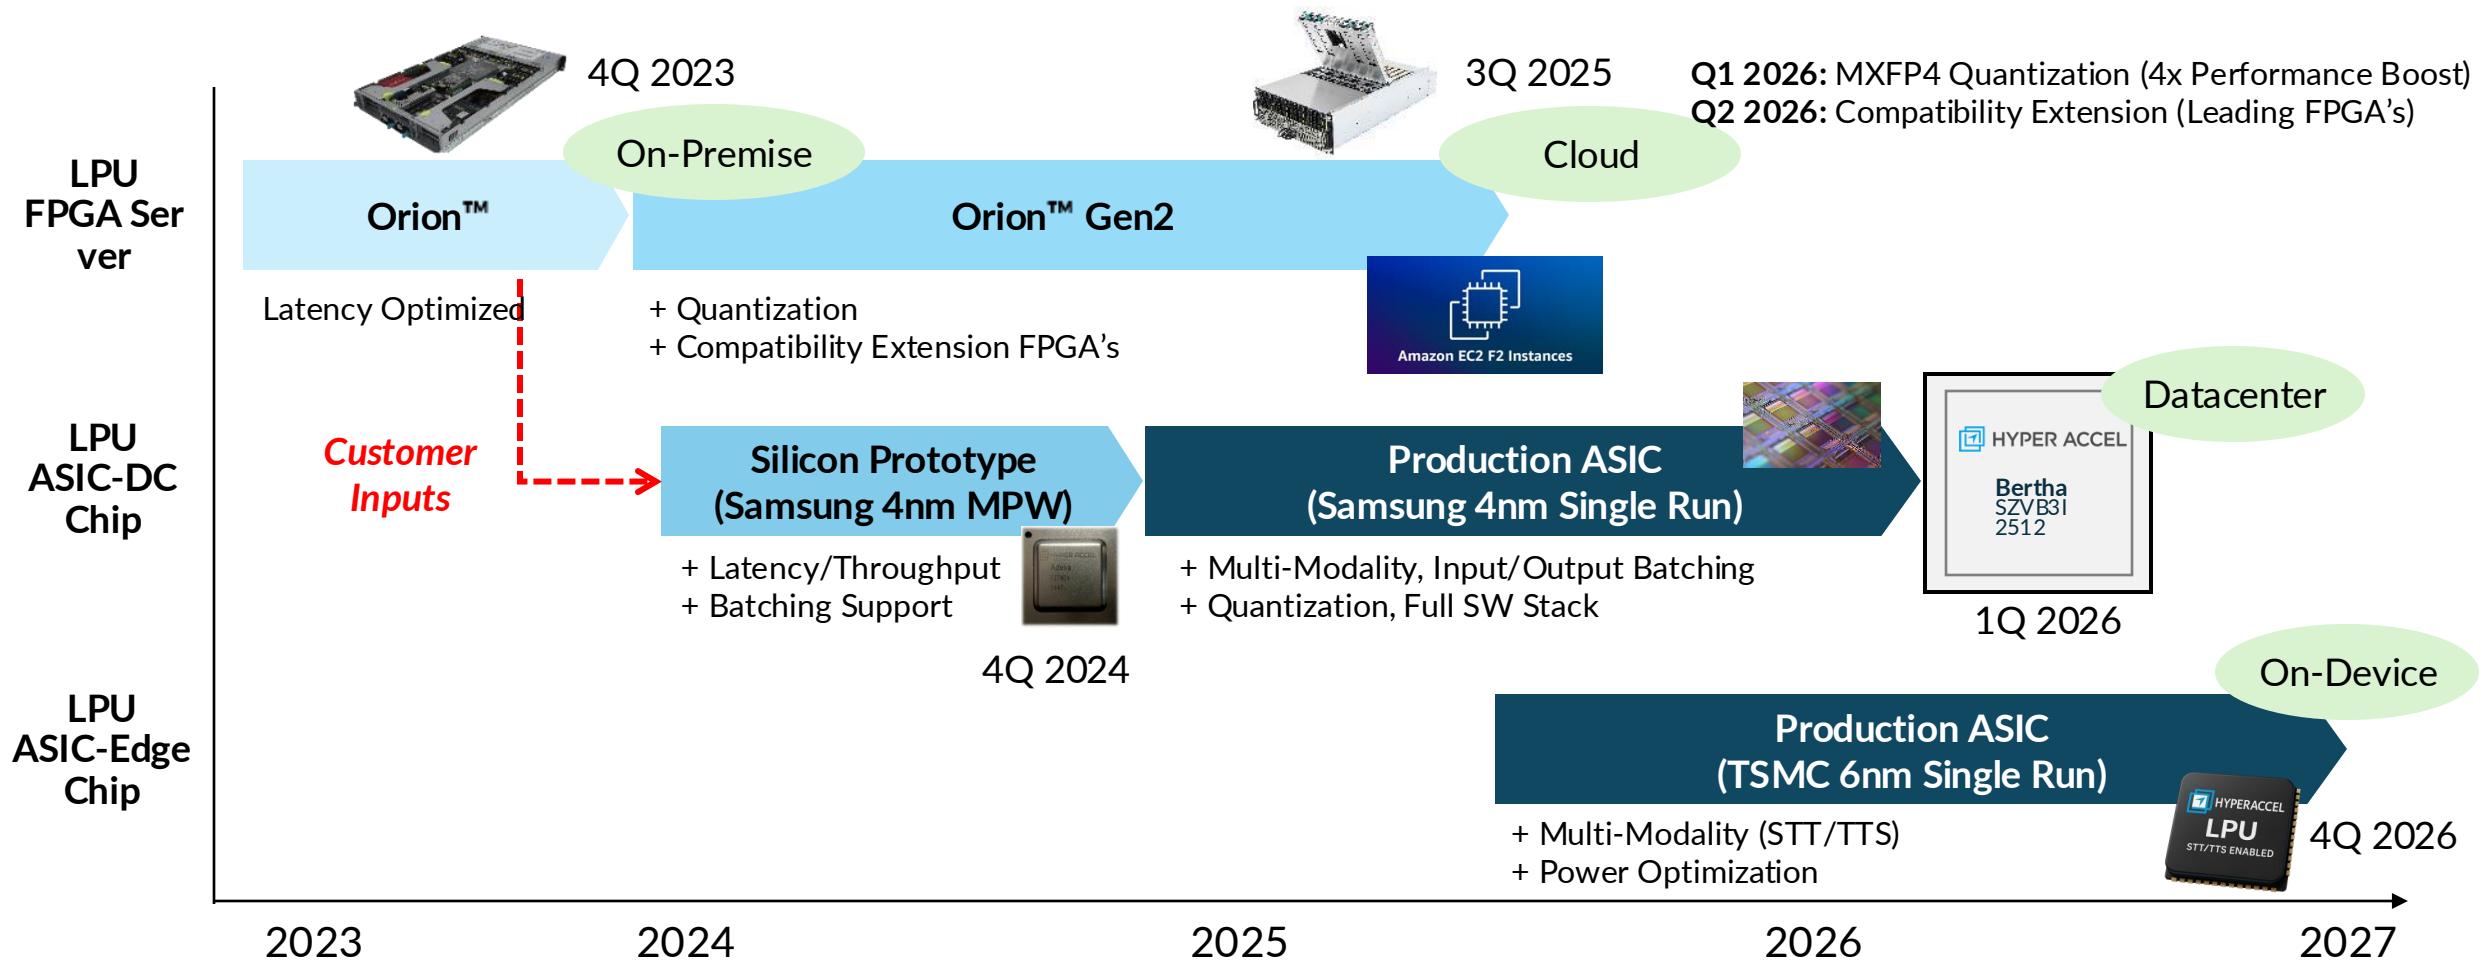


Datacenter Cost Reduction



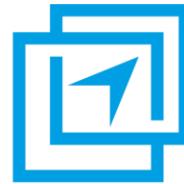
Source: AWS, HyperAccel Analysis

HyperAccel Product Roadmap



HyperAccel Product Portfolio

Product	LPU ASIC-DC Chip Available 1Q 2026	Orion™ LPU FPGA Server Commercially Available	LPU ASIC-Edge Chip Available 4Q 2026
Target Customer	Hyperscalers Cloud Service Providers AI Cloud Providers Datacenters	Telecom Research Labs Universities Startups	Automotive Consumer Electronics/Smart Home AI Devices Robotics IoT
Value Proposition	High Performance Cost & Energy Efficient Attractive Pricing	Affordable Performance Cost & Energy Efficient Realtime Throughput	Edge Performance Cost & Energy Efficient Mass-Market Pricing
Support & Maintenance	HW: Server Partner ASIC: HyperAccel SW: HyperAccel	HW (incl. FPGA): Server Partner FPGA (Image): HyperAccel SW: HyperAccel	HW: OEM Partner ASIC: HyperAccel SW: HyperAccel



HYPER ACCEL

Hyper-Accelerated Solutions For Mission-Critical Workloads